Productizing new technologies for datacenter deployment

Prepared for a Talk at IEEE/ACM International Symposium on Microarchitecture
Cloud@MICRO (10/18/2021)

Siamak Tavallaei
Chief Systems Architect, Google
On the Board of Directors, CXL Consortium
Server Project Lead at OCP Incubation Committee

Points expressed here are not reflective of Google plans or views at-large

Abstract

This section explores driving industry-standard efforts and technologies such as **C**ompute e**X**press **L**ink (**CXL**) into Open Compute Project (**OCP**) as a means to realize products in such a way to enable and encourage high-volume adoption into **datacenters**.

Outline

- Earlier presentations discussed workload performance and a few new technologies
- Customer (workloads & operations)

- CXL as a technology example
  - CXL enables technical benefits
  - CXL-based offerings and use-cases of interest

- Fundamental requirements persist
- Solutions

- Enablers
- Industry-wide efforts: CXL, PCIe, SNIA, NVMe, JEDEC, DMTF, and OCP

A Flywheel of Opportunities for Datacenters

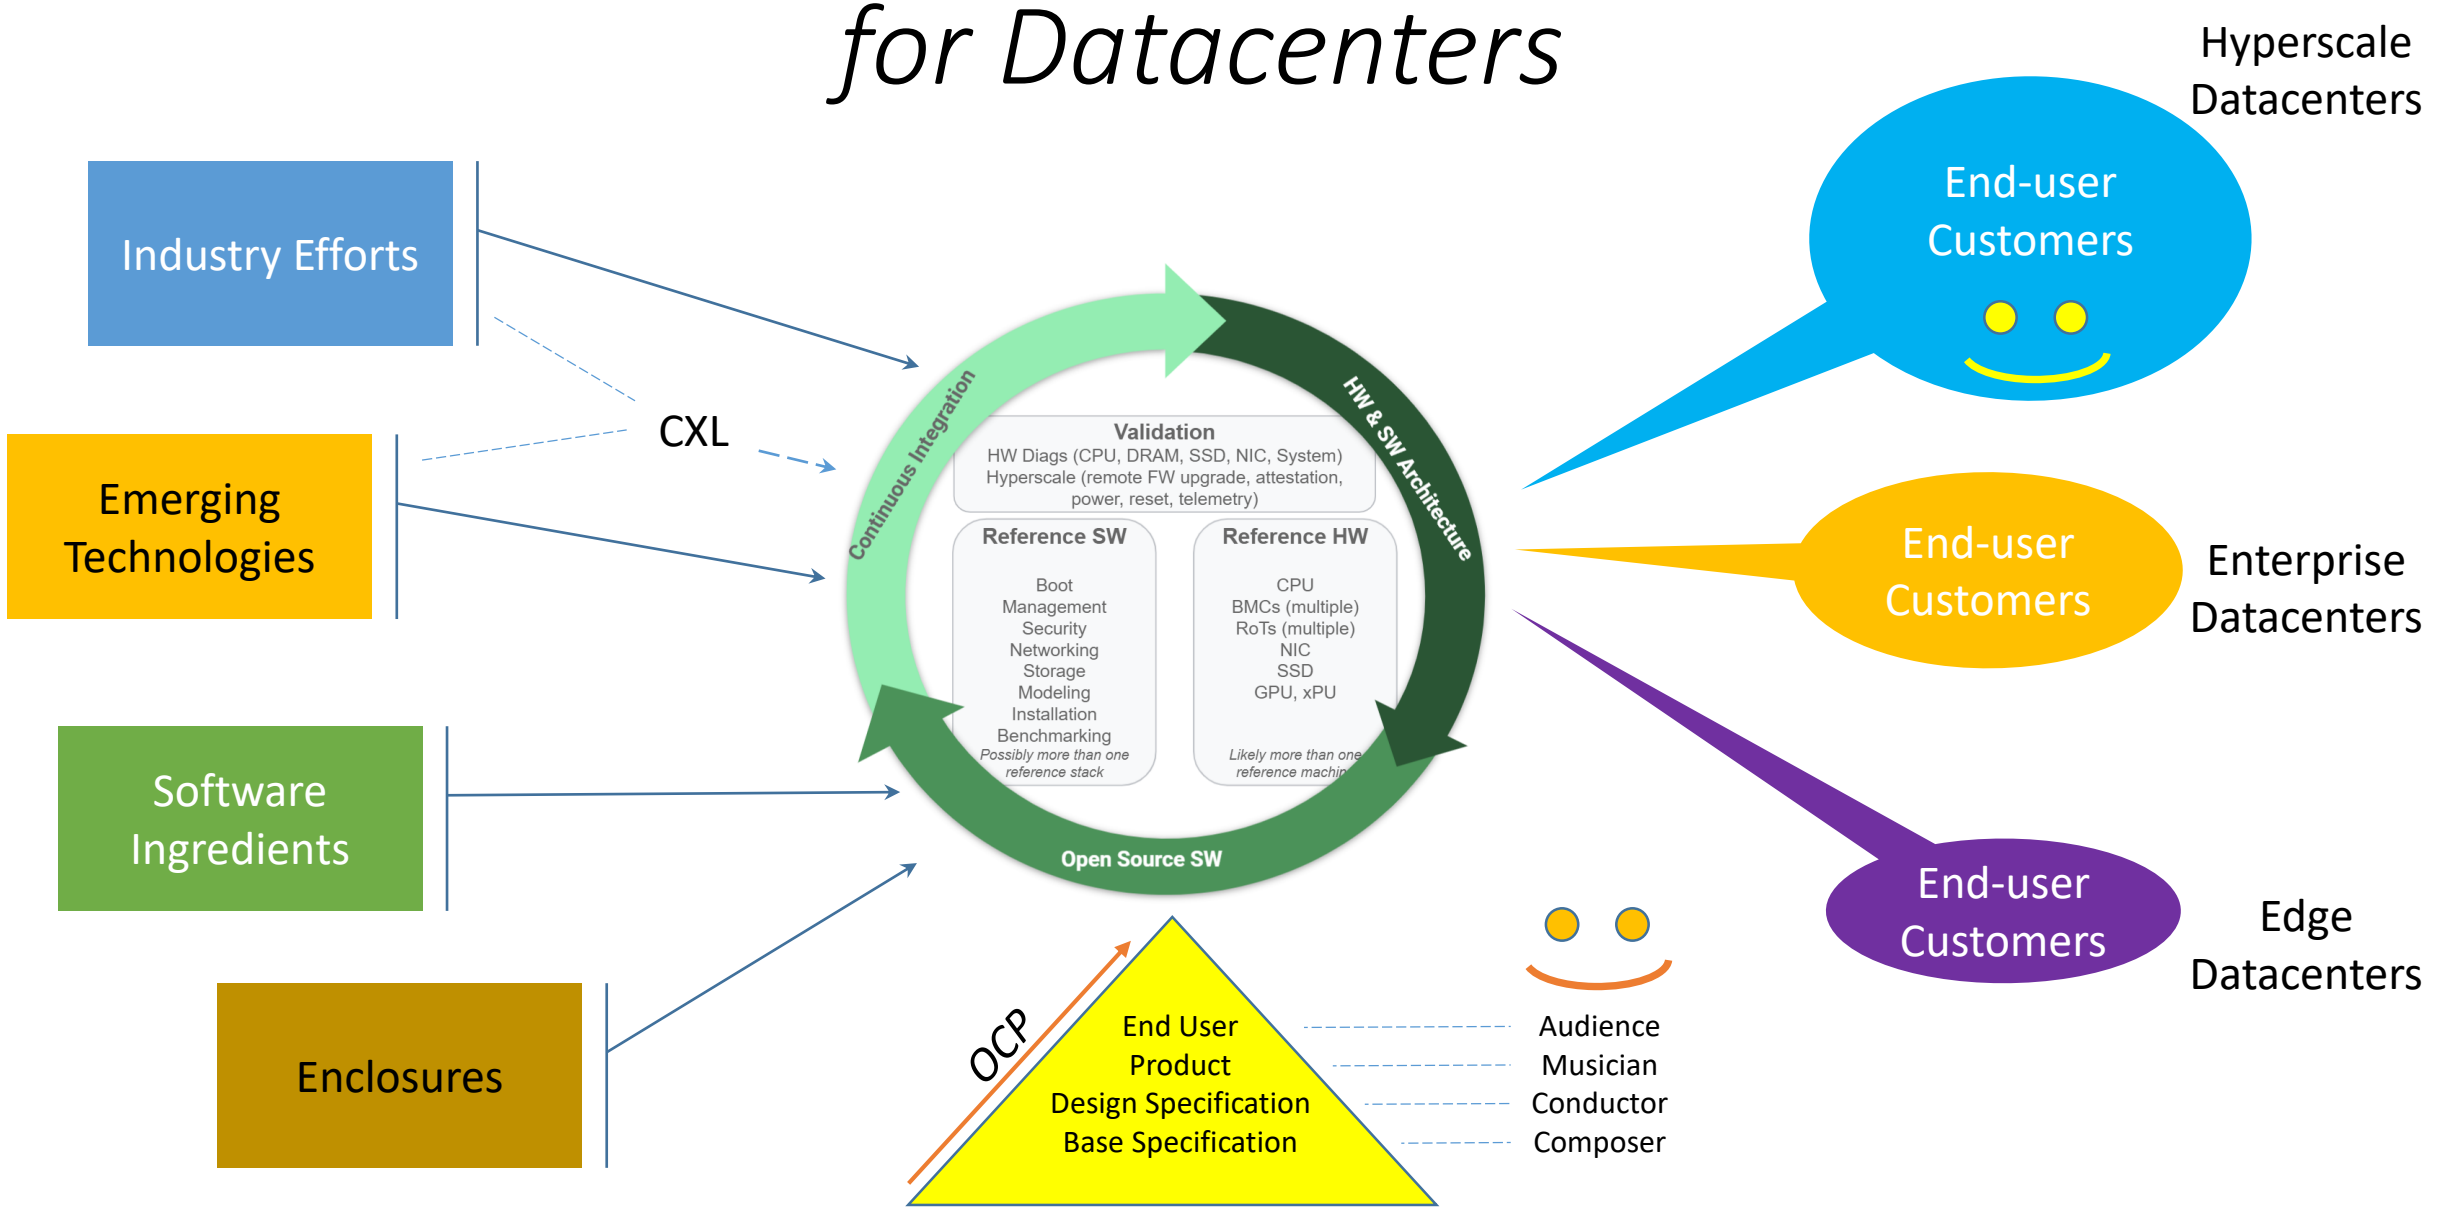

CXL-enabled Opportunities

Interconnect

- Based on PCIe physical layer, **high-speed**
- Optimization for **Coherent**, Load/Store Semantics with **low-latency** for short packets
- Fan-out using a Switch (**large** systems)

Memory

- Memory Capacity and Bandwidth **Expansion**
- Memory **Pooling**
- **Emerging** Memory Technologies

Storage-class Memory

- Architected optimizations for **persistent** memory using Load/Store semantics
- Pooling (**sub-dividing** a Large Device)

Accelerators

- Computational **Off-loading**
- CPU and Accelerator **working on the same coherent memory region**
- **Avoiding** superfluous **data movement** and reducing the associated time and energy (computation applications: in-memory, in-storage, and in-peer-accelerator)

CXL-enabled Opportunities

Interconnect

- Based on PCIe physical layer, **high-speed**
- Optimization for **Coherent**, Load/Store Semantics with **low-latency** for short packets
- Fan-out using a Switch (**large** systems)

Memory

- Memory Capacity and Bandwidth **Expansion**
- Memory **Pooling**
- **Emerging** Memory Technologies

Storage-class Memory

- Architected optimizations for **persistent** memory using Load/Store semantics
- Pooling (**sub-dividing** a Large Device)

Accelerators

- Computational **Off-loading**
- CPU and Accelerator **working on the same coherent memory region**
- **Avoiding** superfluous **data movement** and reducing the associated time and energy (computation applications: in-memory, in-storage, and in-peer-accelerator)

Lots of fun for the technologists!

CXL Ecosystem

Major companies have announced product plans around ***CXL***

- SoC suppliers
- Memory controller suppliers
- Storage suppliers
- Network controller suppliers
- Accelerator suppliers
- Switch suppliers

CXL Ecosystem

Major companies have announced product plans around ***CXL***

- SoC suppliers
- Memory controller suppliers
- Storage suppliers
- Network controller suppliers
- Accelerator suppliers
- Switch suppliers

*This **broad** engagement is the major **advantage** we expect of CXL to deliver a **successful** and profitable environment for all **participants***

Customers

- ***Enterprise*** customers have **diverse** set of needs

- As the *Enterprise* customers move their requirements to the *Cloud Datacenters*, they enjoy the benefits **at-scale** and bring their diverse needs to ***The Cloud***

- ***Edge*** solutions benefit from this Enterprise & Cloud **interplay**

Do off-the-shelf solutions *(**OTS**)*

meet ***Datacenter*** Requirements?

Fundamental Requirements Persist

We still need to deliver integrated **hardware** and **software** solutions which are

Useful

- Desirable

High Quality

- Secure (RoT and Chain of Trust: at-rest, in-transit, and secure execution)
- Safe
- Reliable
- Available

Manageable

- Serviceable
- Diagnosable

Performant

Efficient (power, space, cost, time, complexity, …)

Fundamental Requirements Persist

We still need to deliver integrated **hardware** and **software** solutions which are

- Useful
  - Desirable
- High Quality
  - Secure (RoT and Chain of Trust: at-rest, in-transit, and secure execution)
  - Safe
  - Reliable
  - Available
- Manageable
  - Serviceable
  - Diagnosable
- Performant
- Efficient (power, space, cost, time, complexity, …)

*Especially when driving the solutions into **Large Datacenters***

Solution

Balanced **Core** Architecture

- Frameworks, Software, Compute, BW, Capacity, Latency

General-purpose

- Modular **Building Blocks**

Extensible

- Allow heterogeneous variants based on the core Building Blocks

Challenges in deploying traditional Servers into Hyperscaled Datacenters

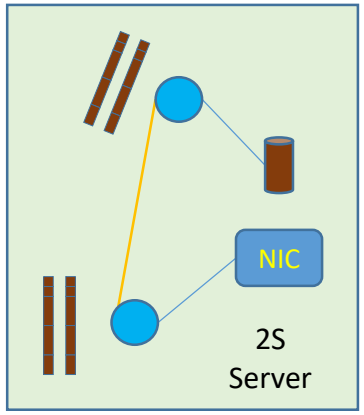

Balanced match of:

- CPU core count
- Memory Capacity, Bandwidth, and Latency
- Storage Capacity, Bandwidth, IOPS, and Tail Latency
- Network Bandwidth

Challenging to meet the above balance in presence of varied workloads and customer VMs in high-volume production

Result:

- **Overprovisioned** resources to meet customer demands
- Unused or underused resources
- Increased Cost

How can a new technology

such as CXL

help?

Extensible Solutions

Topology
- Point-to-point
- Multi-port
- Switched

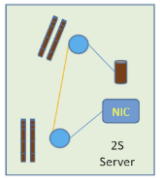

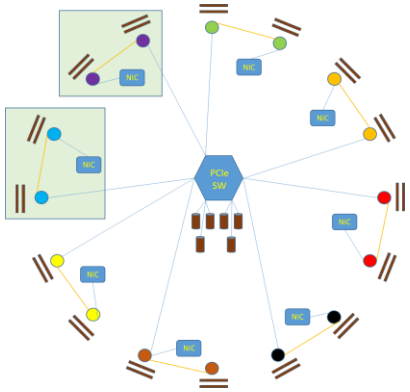

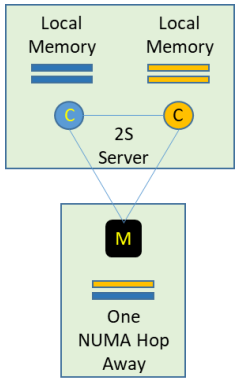

Density (multi-port)
- Dense packaging of (n x m) multi-ported Devices
- Liquid cooling

Reach (SERDES)
- Longer Links *(to all devices including memory!)*
- Modular Enclosures
- Cabled Solutions
- Photonics

Extensibility (heterogeneous)
- Compute (xPU), Memory, Storage, Networking

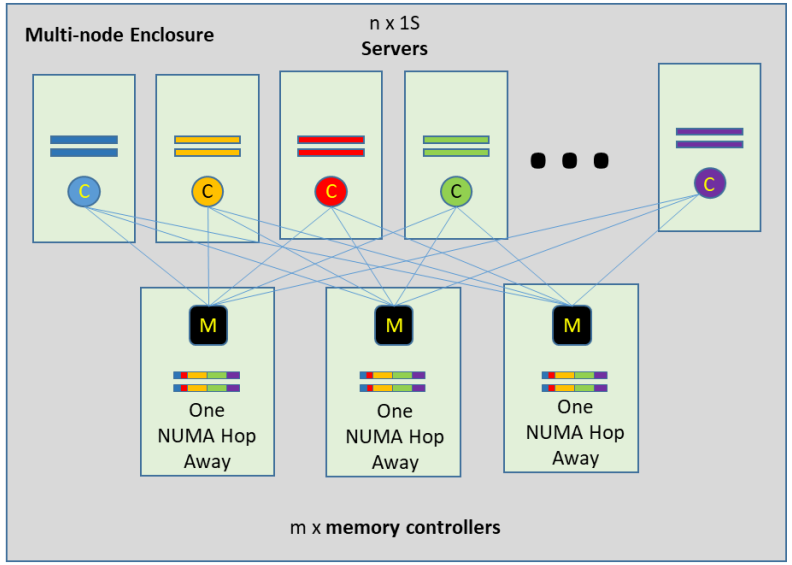

Capacity-matched
Bandwidth-matched

Is CXL the End-all?

Should we move everything
to the new technology?

Is CXL the End-all?

Should we move everything
*to the **new technology**?*

Putting Things where they Belong!

PCIe

- Software-managed consistency (DMA, RDMA)
- Block Data Moves (Large Payloads)
- Deferred Calls, Interrupts
- Latency-tolerant
- Sequential data access

CXL

- Hardware-managed (Coherence)
- Load/Store (Short Packets)
- In-line codes
- Latency-sensitive
- Concurrent data access

Putting Things where they Belong!

PCIe

- Software-managed consistency (DMA, RDMA)
- Block Data Moves (Large Payloads)
- Deferred Calls, Interrupts
- Latency-tolerant
- Sequential data access

CXL

- Hardware-managed (Coherence)
- Load/Store (Short Packets)
- In-line codes
- Latency-sensitive
- Concurrent data access

- Enabling new optimizations and programing paradigms

There will be a transitional period

from one
to the other

Remember!

All Fundamental Requirements Persist!

Enablers Needed

Enablers Needed

A Vibrant Ecosystem

Enablers

Riding on the Coattail of The Giants

Industry Efforts

- CXL Consortium
- PCIe SIG
- SNIA
- NVMe
- JEDEC
- DMTF
- OCP
- …

Enablers (Software and Firmware Ingredients)

CXL Fabric Manager

- Secure composability, allocation, on-lining/off-lining

Pre-boot Environment

- Discovery, Enumeration

CXL Bus Driver

- Configuration, Resource allocation

CXL Memory Device Driver

- Interactions with Bus Driver, Fabric Manager, and VMM
- RAS, Security, Fault-isolation, On-lining, Off-lining, …

  ECN: Error Isolation on CXL.mem and CXL.cache (Enabled by the Root Port; requires Software Stack to recover from faults)

OS-specific Software

- VMM, Hypervisor
- VM Allocation, Orchestration, Fault-isolation & Recovery

Datacenter-ready Integrated System *(DC-Stack)*

CXL Consortium relies on other standards bodies to provide specifications for suitable mechanical/thermal/electrical **Form Factors**

Newly enabled solutions based on CXL will require new form factors

Reference Designs (open-source hardware at OCP)

PoCs

Datacenter-ready Integrated System *(DC-Stack)*

CXL Consortium relies on other standards bodies to provide specifications for suitable mechanical/thermal/electrical **Form Factors**

Newly enabled solutions based on CXL will require new form factors

Reference Designs (open-source hardware at OCP)

PoCs

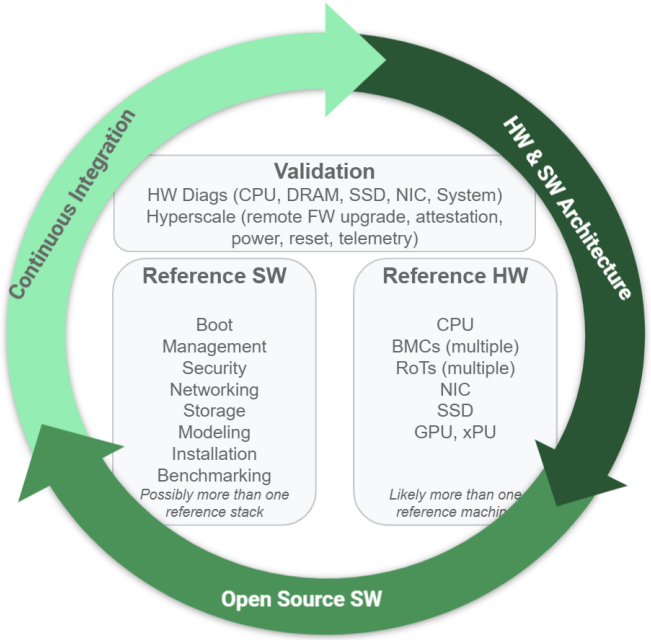

Datacenter-ready Integrated System *(DC-Stack)*

A convolution of many essential ingredients

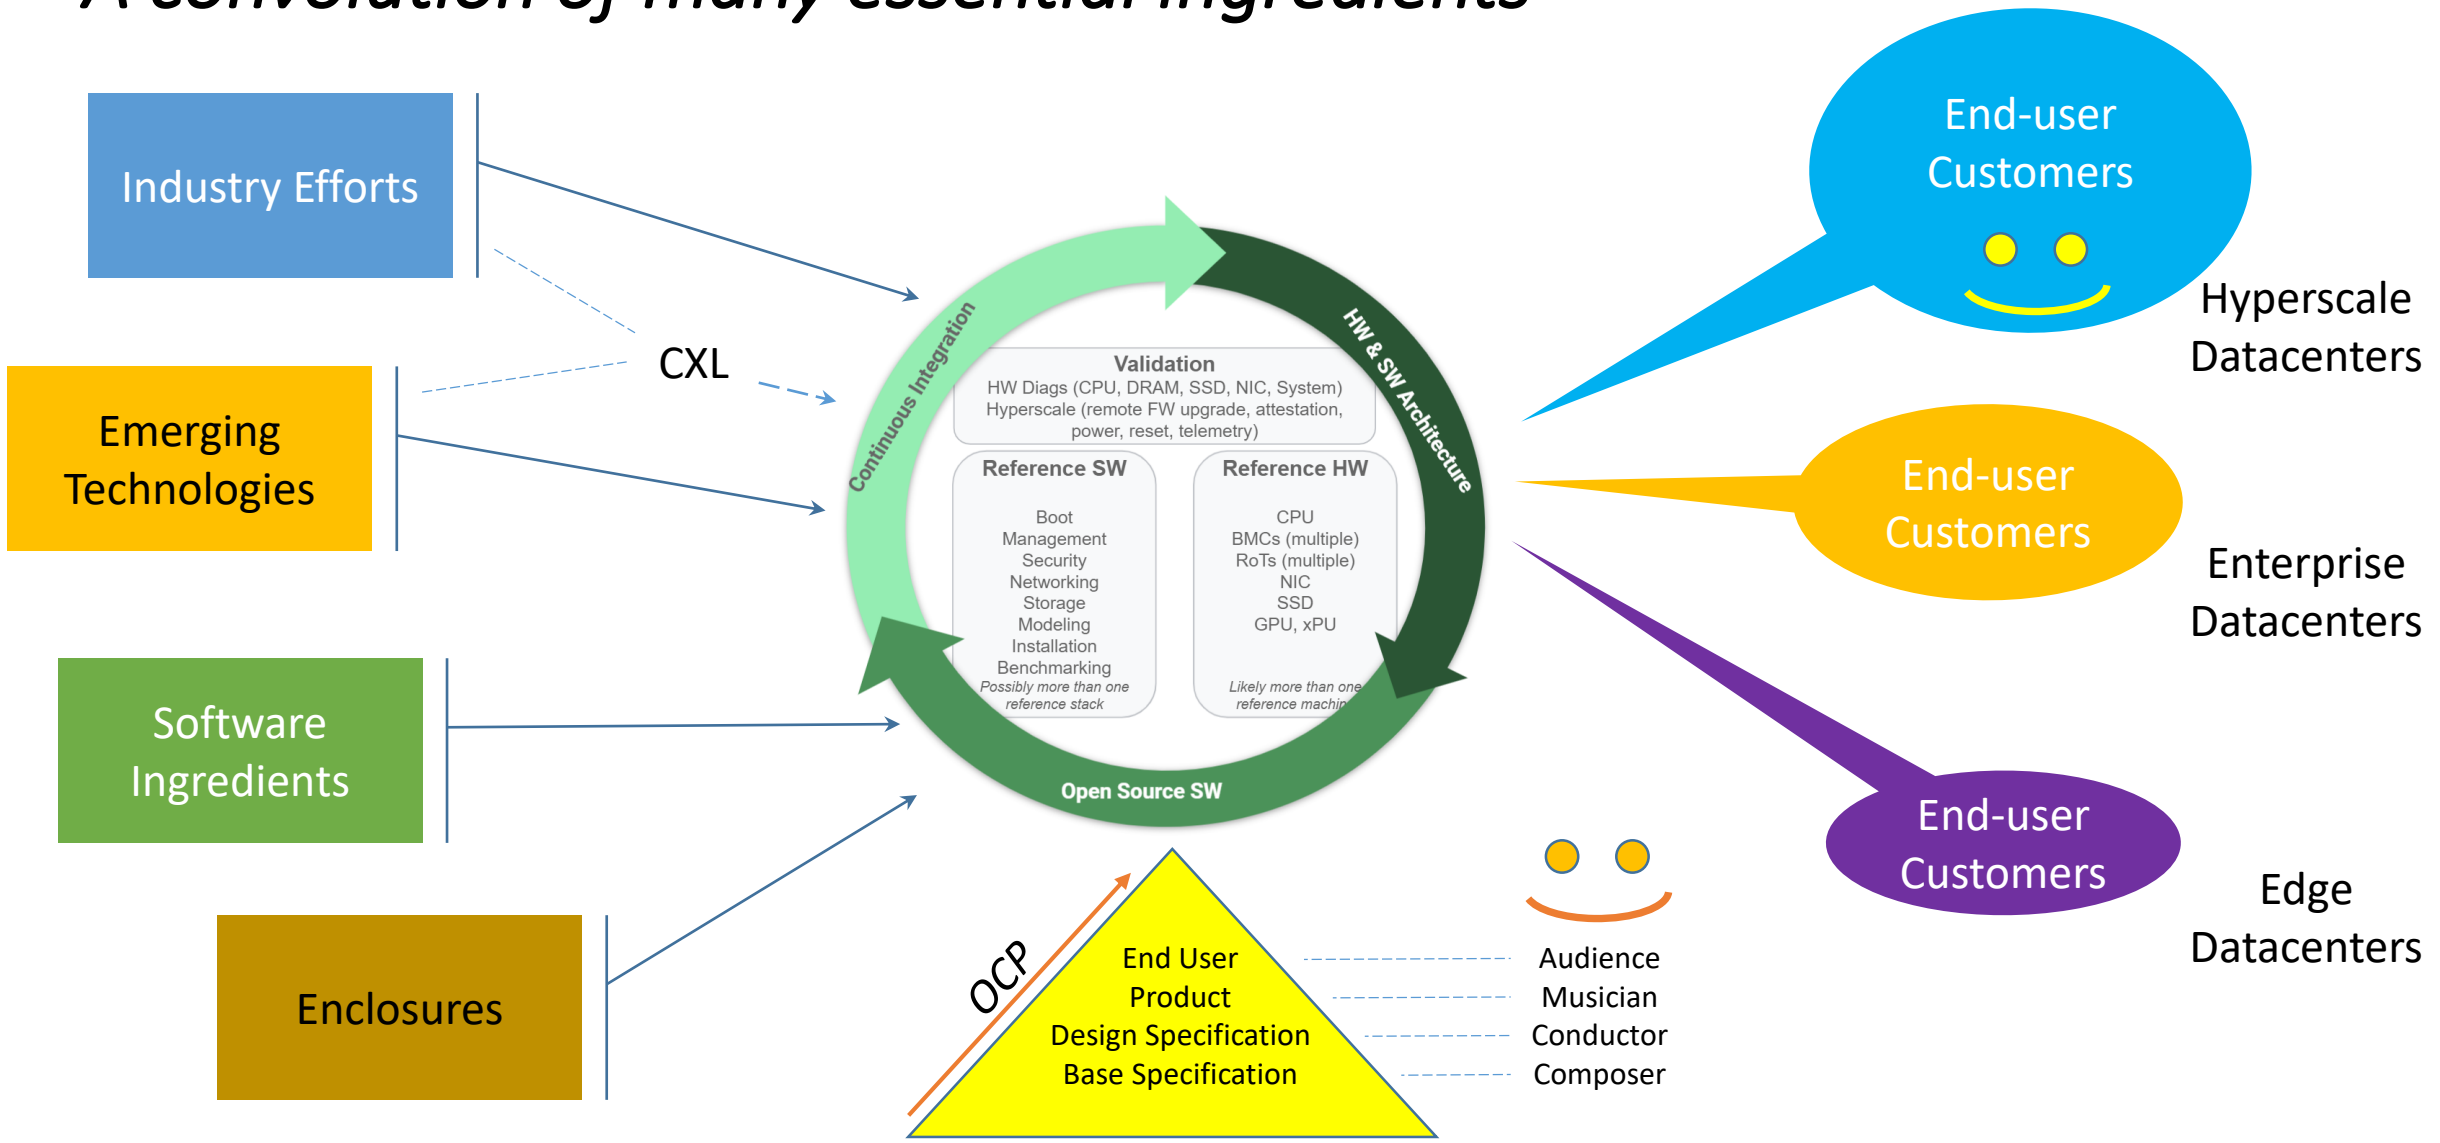

For OCP Discussions

Modular Hardware

Reference designs for three different architectural instances

1S/2S Server

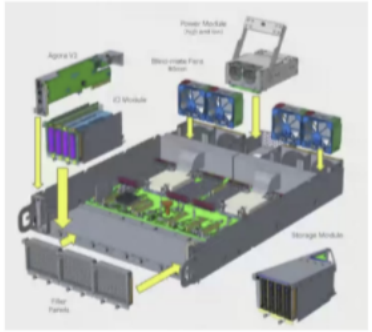

Rack
Power
Cooling
Multi-Blade/Instance Chassis
HPM (CPUs + Memory) per Blade
Multiple DC-SCMs or Multi-Host DC-SCM
DC-MIO
IO Module/Cage (IO Slots)
Multiple SmartNICs or Multi-Host SmartNIC

Multi-CPU Server

Rack
Power
Cooling
Single Instance Chassis
Multi-HPM (CPUs + Memory)
IO Module/Cage (IO Slots)
Single-Host DC-SCMS
DC-MIO
Single-host SmartNIC

GPU/Acc. or Storage Expansion Chassis

Expansion Chassis, Power, Cooling
Head Node (Server)
Interconnect (cables, retimer?)
Out-of-band Management
BMC (or a variant of DC-SCM)

The Hardware Portion

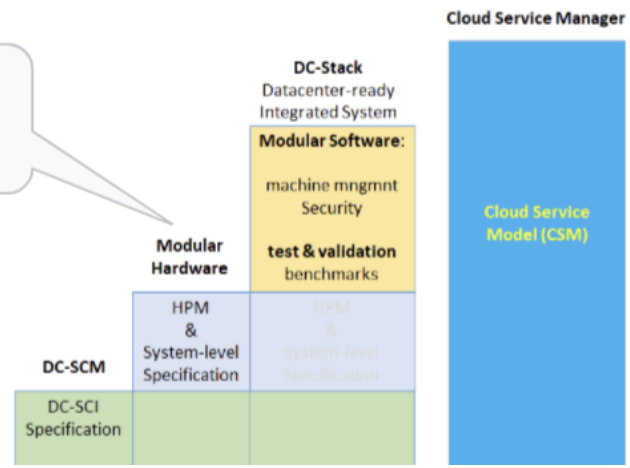

Summary

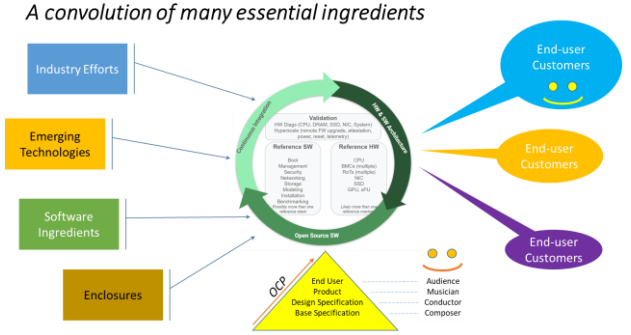

- A new technology such as CXL enables special **benefits**, but we still need to deliver the **fundamental requirements**
- Delivering these technical advantages will take major **ecosystem** effort from various industry players
- ***You*** along with balanced, extensible, **modular** solutions, along with the staged software stack are the **enablers**
- Taking advantage of the industry-wide efforts, we can deliver CXL-based **PoCs** toward a datacenter-ready integrated system ***(DC-Stack)*** via open-sourced hardware and software ***(OCP)***

You *are a*

Giant ***Enabler!***

Lend your Coattail!

Available presentations on ***C****ompute e****X****press* ***L****ink (CXL)* as an open-standard specification

The material that CXL has published on memory pooling:

Webinar: Compute Express Link™ 2.0 Specification: Memory Pooling

LinkedIn Post on CXL Memory Pooling

Recap Q&A Blog: Part 1

Recap Q&A Blog: Part 2

CXL Memory Pooling Slides

CXL Memory Pooling Video

CXL 2.0 Animated Video

CXL to Gen-Z Use Cases

Presentations on Datacenter-ready Integrated System **(DC-Stack):**

DC-Stack_Datacenter-ready Integrated System_OCP

http://files.opencompute.org/oc/public.php?service=files&t=dd1e012f85ab59a608d758db8357539c